



Deploying JupyterHub
for students and
researchers

JupyterCon 2017

Min Ragan-Kelley, Simula
Carol Willing, Cal Poly
Yuvi Panda, UC Berkeley
Ryan Lovett, UC Berkeley

Set Up

```
git clone \
https://github.com/jupyterhub/jupyterhub-tutorial \
/srv/jupyterhub
```



Tutorial logistics

- 9:00 - 9:15 Welcome
- 9:15 - 10:30 JupyterHub
- 10:30 - 11:00 Morning break
- 11:00 - 12:20 JupyterHub and Kubernetes
- 12:20 - 12:30 Wrap up



Tutorial logistics

- Gitter channel for this tutorial

<https://gitter.im/jupyterhub-tutorial/Lobby>

- Post-its

questions and assistance

done with task and doing well

- Tutorial materials

<https://github.com/jupyterhub/jupyterhub-tutorial>



What are Jupyter and JupyterHub?



What is a Notebook?

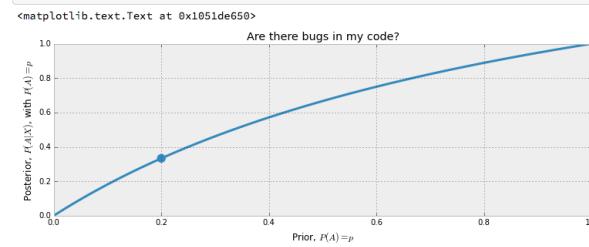
- Document
- Environment
- Web app

We have already computed $P(X|A)$ above. On the other hand, $P(X \sim A)$ is subjective: our code can pass tests but still have a bug in it, though the probability there is a bug present is reduced. Note this is dependent on the number of tests performed, the degree of complication in the tests, etc. Let's be conservative and assign $P(X \sim A) = 0.5$. Then

$$P(A|X) = \frac{1 \cdot p}{1 \cdot p + 0.5(1 - p)}$$
$$= \frac{2p}{1 + p}$$

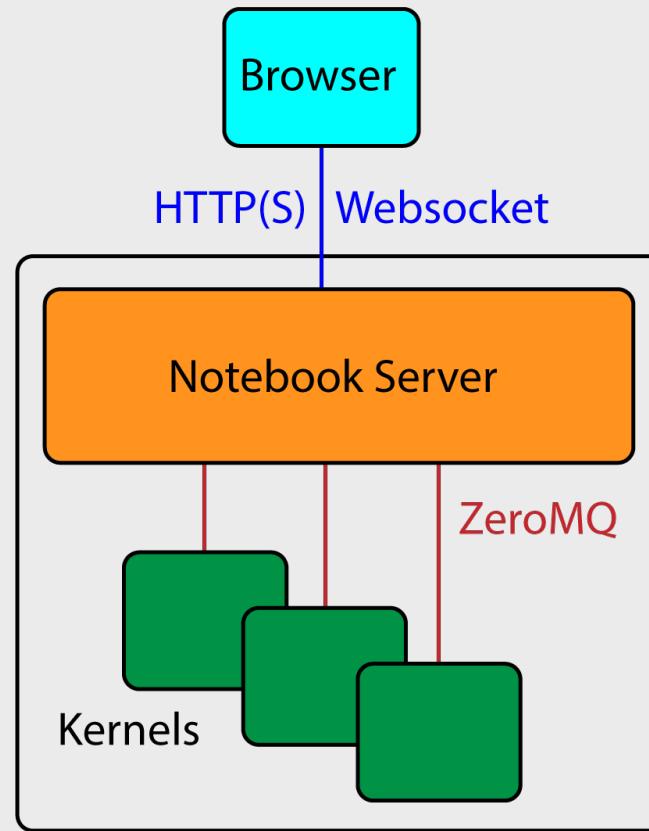
This is the posterior probability. What does it look like as a function of our prior, $p \in [0, 1]$?

```
figsize(12.5, 4)
p = np.linspace(0, 1, 50)
plt.plot(p, 2 * p / (1 + p), color="#348ABD", lw=3)
# plt.fill_between(p, 2*p/(1+p), alpha=.5, facecolor="#A60628")
plt.sca((0.2, 2 * (0.2) / 1.2, 0, 1.2, s=140, c="#348ABD")
plt.ylim(0, 1)
plt.xlabel("Prior, $P(A) = p$")
plt.ylabel("Posterior, $P(A|X)$, with $P(A) = p$")
plt.title("Are there bugs in my code?")
```



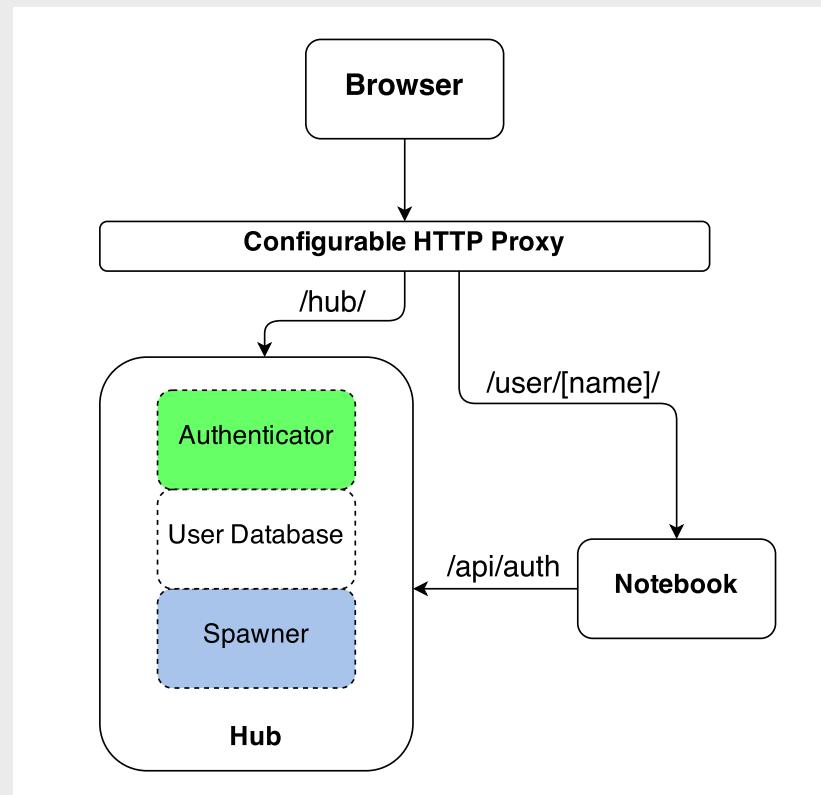
<https://github.com/CamDavidsonPilon/Probabilistic-Programming-and-Bayesian-Methods-for-Hackers>

What is a Notebook Server?



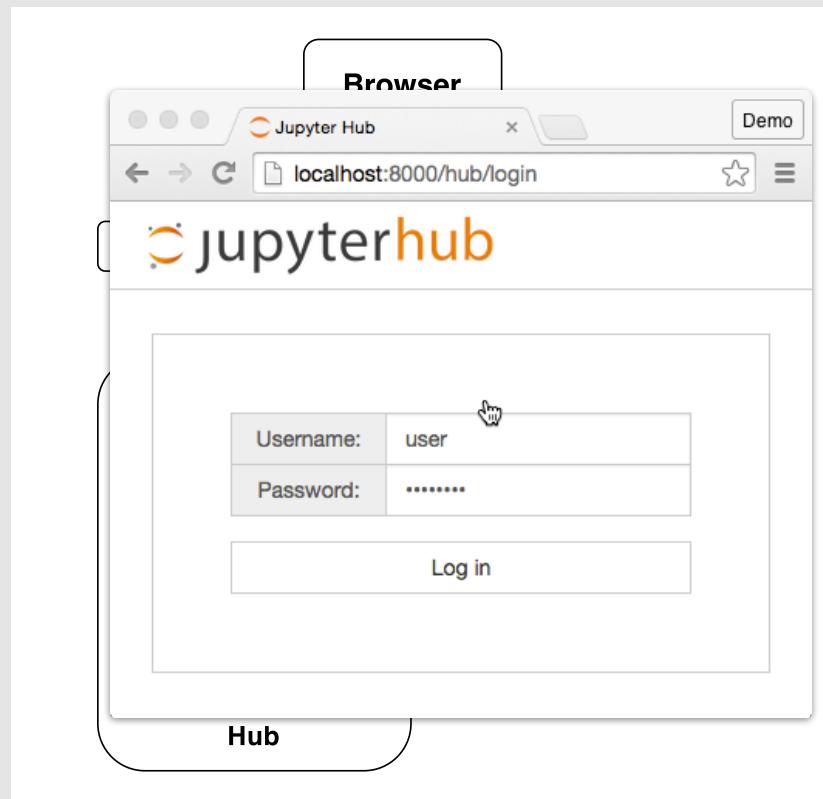
jupyterhub

- Manages authentication
- Spawns single-user servers on-demand
- Each user gets a complete notebook server



jupyterhub

- Initial request is handled by Hub
- User authenticates via form / OAuth
- Spawner starts single-user server
- Hub notifies Proxy
- Redirects user to /user/[name]
- Single-user Server verifies auth with Hub



Installation



Photo taken by Matthew Bowers

Installation (as administrator)

conda:

```
conda install -c conda-forge jupyterhub  
conda install notebook
```

pip, npm:

```
python3 -m pip install jupyterhub  
npm install -g configurable-http-proxy
```

test:

```
jupyterhub -h  
configurable-http-proxy -h
```



Installation (this repo)

```
conda env create -f environment.yml  
source activate jupyterhub-tutorial
```



Installation Caveats

JupyterHub installation must be
readable+executable by all users*

This is often not the case for envs, so be
careful

*when using local users



Plug: conda-forge



Community-managed conda packages.

<https://conda-forge.github.io>

`conda config --add channels conda-forge`

Installation: Spawner and Single User

<https://docs.docker.com/engine/installation>

```
pip install dockerspawner  
docker pull jupyterhub/singleuser
```



JupyterHub Defaults

- Authentication: PAM (local users, passwords)
- Spawning: Local users
- Hub must run as root



Aside: SSL

- JupyterHub is an authenticated service - users login.
That should **never** happen over plain HTTP.
- For testing, we can generate self-signed certificates:

```
openssl req -x509 -nodes -days 365 -newkey rsa:1024 \  
-keyout jupyterhub.key -out jupyterhub.crt
```

Note: Safari will not connect websockets to untrusted (self-signed) certs



SSL: Let's Encrypt

- <https://letsencrypt.org/getting-started/>
- Free SSL for any domain

```
wget https://dl.eff.org/certbot-auto
chmod a+x certbot-auto
./certbot-auto certonly --standalone -d mydomain.tld
```

```
key: /etc/letsencrypt/live/mydomain.tld/privkey.pem
cert: /etc/letsencrypt/live/mydomain.tld/fullchain.pem
```



Configuration



Photo taken by Matthew Bowers

Start configuring JupyterHub

```
jupyterhub --generate-config
```

```
c.JupyterHub.ssl_key = 'jupyterhub.key'  
c.JupyterHub.ssl_cert = 'jupyterhub.crt'  
c.JupyterHub.port = 443
```

jupyterhub_config.py



Installing language kernels for all users

```
conda create -n py2 python=2 ipykernel  
conda run -n py2 -- ipython kernel install  
  
jupyter kernelspec list
```



Using GitHub OAuth

<https://github.com/settings/applications/new>

Register a new OAuth application

Application name

Something users will recognize and trust

Homepage URL

The full URL to your application homepage

Application description

This is displayed to all potential users of your application

Authorization callback URL

Your application's callback URL. Read our [OAuth documentation](#) for more information.

[Register application](#)

JupyterHub tutorial

 **minrk** owns this application. [Transfer ownership](#).

1 user

Client ID
862e21790aeb654a0c98

Client Secret
c2c1 [REDACTED] befa6

[Revoke all user tokens](#) [Reset client secret](#)

Using GitHub OAuth

File: ./env

```
export GITHUB_CLIENT_ID=from_github  
export GITHUB_CLIENT_SECRET=from_github  
export OAUTH_CALLBACK_URL=https://YOURDOMAIN/hub/oauth_callback
```

```
source ./env
```



Using GitHub OAuth

We need to install: OAuthenticator

```
python3 -m pip install oauthenticator
```

Config file: jupyterhub_config.py

```
from oauthenticator.github import LocalGitHubOAuthenticator
c.JupyterHub.authenticator_class = LocalGitHubOAuthenticator
c.LocalGitHubOAuthenticator.create_system_users = True
```



Specifying Users

By default, any user that successfully authenticates is allowed to use the Hub.

This is appropriate for shared workstations with PAM Auth, but probably not GitHub:

```
# set of users allowed to use the Hub
c.Authenticator.whitelist = {'minrk', 'takluyver'}
```



```
# set of users who can administer the Hub itself
c.Authenticator.admin_users = {'minrk'}
```



Custom Authenticators



Photo taken by Matthew Bowers



Using DockerSpawner

We need DockerSpawner:

```
python3 -m pip install dockerspawner netifaces  
docker pull jupyterhub/singleuser
```

In jupyterhub_config.py:

```
from oauthenticator.github import GitHubOAuthenticator  
c.JupyterHub.authenticator_class = GitHubOAuthenticator  
  
from dockerspawner import DockerSpawner  
c.JupyterHub.spawner_class = DockerSpawner
```



Using DockerSpawner

```
from dockerspawner import DockerSpawner
c.JupyterHub.spawner_class = DockerSpawner

# The Hub's API listens on localhost by default,
# but docker containers can't see that.
# Tell the Hub to listen on its docker network:
import netifaces
docker0 = netifaces.ifaddresses('docker0')
docker0_ipv4 = docker0[netifaces.AF_INET][0]
c.JupyterHub.hub_ip = docker0_ipv4['addr']
```



Using DockerSpawner

- There is *loads* to configure with Docker
- Networking configuration
- Data volumes
- `DockerSpawner.image = 'jupyter/scipy-notebook:8f56e3c47fec'`



Customizing Spawners



Photo taken by Matthew Bowers



JupyterHub with supervisor

apt-get install supervisor

```
#!/usr/bin/env bash
# /srv/jupyterhub/launch.sh
set -e
source ./env
exec jupyterhub $@
```

```
# /etc/supervisor/conf.d/jupyterhub.conf
[program:jupyterhub]
command=bash launch.sh
directory=/srv/jupyterhub
autostart=true
autorestart=true
startretries=3
exitcodes=0,2
stopsignal=TERM
redirect_stderr=true
stdout_logfile=/var/log/jupyterhub.log
stdout_logfile_maxbytes=1MB
stdout_logfile_backups=10
user=root
```

Reference Deployments

<https://github.com/jupyterhub/helm-chart>

Helm, KubeSpawner, Hub in container

<https://github.com/jupyterhub/jupyterhub-deploy-docker>

docker-compose, DockerSpawner, Hub in Docker

<https://github.com/jupyterhub/jupyterhub-deploy-teaching>

ansible, no docker, nbgrader



Docker Deployment

- Docker Compose: <https://docs.docker.com/compose/install/>
- git clone <https://github.com/jupyterhub/jupyterhub-deploy-docker>
- Setup the basics (creates volumes, network, pulls images):
`make build`



Docker Deployment

- `mkdir secrets`
- Copy SSL key, cert to:
 - `secrets/jupyterhub.cert` (cert)
 - `secrets/jupyterhub.key` (key)



Docker Deployment

Make userlist:

minrk admin
willingc admin
yuvipanda
ryanlovett



Docker Deployment

Launch: 

`docker-compose up`



Optimizations and Best Practices



Photo taken by Matthew Bowers



Optimizations and best practices

- **Always use SSL!**
- Use postgres for the Hub database
- Put nginx in front of the proxy
- Run cull-idle-servers service to prune resources
- Global configuration in /etc/jupyter and /etc/ipython
- **Back up your user data!!!**



When to use JupyterHub

- A class where students can do homework (nbgrader)
- A short-lived workshop, especially if installation is hard
- A research group with a shared workstation or small cluster
- On-site computing resources for researchers and analysts at an institution



When ***not*** to use JupyterHub

- JupyterHub is Authenticated and Persistent
- It takes work to keep it going
- SageMathCloud is *hosted* and provides realtime-collaboration
- A growing number of cloud providers support notebook services (Google, Microsoft, etc.)



JupyterHub API



Photo taken by Matthew Bowers



End of JupyterHub Fundamentals

Break - 30 min
See you back here at 11:00 am.

Up next:
JupyterHub and
Kubernetes



Attributions and recognition

A huge thank you to the Project Jupyter team and community.
Your hard work and passion makes this all possible.





Thank you



Thank you

try.jupyter.org

www.jupyter.org

ipython.org

numfocus.org

